

Q1

The file valuation.csv (available on Brightspace¹) contains market historical data of real estate valuation collected from Sindian Dist., New Taipei City, Taiwan. Our goal is to predict the per-unit-area house price using several important factors that can potentially affect buyers' decision-making.

Features in the dataset:

X1: the transaction date (for example, 2013.250=2013 March, 2013.500=2013 June, etc.)
X2: the house age (unit: year)
X3: the distance to the nearest MRT station (unit: meter)
X4: the number of convenience stores in the living circle on foot (integer)
X5: the geographic coordinate, latitude. (unit: degree)
X6: the geographic coordinate, longitude. (unit: degree)
Y: house price of unit area (10000 New Taiwan Dollar/Ping, where Ping is a local unit, 1 Ping = 3.3 square meters)

- a. Load the dataset into R from valuation.csv. Split the dataset into training data (first 200 rows) and test data (the rest rows). Name the training and test data valuation_train and valuation_test, respectively. Where is the 11th house in the test data located (answer with the latitude and longitude)?

Solution:

```
# Q1.a
#Load dataset into R from 'valuation.csv'
valuation <- read.csv("C:/Users/Lizi/Downloads/valuation.csv")

#Split the dataset into training data and test data
valuation_train <- valuation[1:200, ]
valuation_test <- valuation[-(1:200), ]

#Check to see if observations line up with excel
head(valuation_train)

##   No      X1    X2      X3  X4      X5    X6     Y
## 1  1 2012.917 32.0  84.87882 10 24.98298 121.5402 37.9
## 2  2 2012.917 19.5 306.59470  9 24.98034 121.5395 42.2
## 3  3 2013.583 13.3 561.98450  5 24.98746 121.5439 47.3
## 4  4 2013.500 13.3 561.98450  5 24.98746 121.5439 54.8
## 5  5 2012.833  5.0 390.56840  5 24.97937 121.5425 43.1
## 6  6 2012.667  7.1 2175.03000  3 24.96305 121.5125 32.1
```

¹ Original source <https://archive.ics.uci.edu/dataset/477/real+estate+valuation+data+set>

```

tail(valuation_train)

##      No      X1      X2      X3  X4      X5      X6    Y
## 195 195 2013.500 15.2 3771.8950 0 24.93363 121.5116 29.3
## 196 196 2013.333 15.2 461.1016 5 24.95425 121.5399 34.6
## 197 197 2013.000 22.8 707.9067 2 24.98100 121.5471 36.6
## 198 198 2013.250 34.4 126.7286 8 24.96881 121.5409 48.2
## 199 199 2013.083 34.0 157.6052 7 24.96628 121.5420 39.1
## 200 200 2013.417 18.2 451.6419 8 24.96945 121.5449 31.6

head(valuation_test)

##      No      X1      X2      X3  X4      X5      X6    Y
## 201 201 2013.417 17.4 995.7554 0 24.96305 121.5491 25.5
## 202 202 2013.417 13.1 561.9845 5 24.98746 121.5439 45.9
## 203 203 2012.917 38.3 642.6985 3 24.97559 121.5371 31.5
## 204 204 2012.667 15.6 289.3248 5 24.98203 121.5435 46.1
## 205 205 2013.000 18.0 1414.8370 1 24.95182 121.5489 26.6
## 206 206 2013.083 12.8 1449.7220 3 24.97289 121.5173 21.4

tail(valuation_test)

##      No      X1      X2      X3  X4      X5      X6    Y
## 409 409 2013.417 18.5 2175.74400 3 24.96330 121.5124 28.1
## 410 410 2013.000 13.7 4082.01500 0 24.94155 121.5038 15.4
## 411 411 2012.667  5.6  90.45606 9 24.97433 121.5431 50.0
## 412 412 2013.250 18.8 390.96960 7 24.97923 121.5399 40.6
## 413 413 2013.000   8.1 104.81010 5 24.96674 121.5407 52.5
## 414 414 2013.500   6.5  90.45606 9 24.97433 121.5431 63.9

#11th house Location in test data
valuation_test[11, c("X5","X6")]

##          X5          X6
## 211 24.97937 121.5425

```

- The **11th house** in the test data is located at **latitude, longitude = (24.97937, 121.5425)**.
- b. Use the training data `valuation_train` to fit a simple linear regression model of Y on $X2$. Interpret the coefficients β_0 (intercept) and β_1 (slope) in terms of the real context. Compute R^2 , training MSE and test MSE.

Solution:

```

#Q1.b
#Fit a simple linear regression of model 'Y' on 'X2'
Y_X2 <- lm(Y~X2, data = valuation_train)

#interpret the coefficients Beta0 and Beta1
summary(Y_X2)

```

```

## 
## Call:
## lm(formula = Y ~ X2, data = valuation_train)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -31.859 -11.634   2.028   8.936  31.496
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 44.46714   1.70482  26.083 < 2e-16 ***
## X2          -0.33841   0.08021  -4.219 3.73e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.97 on 198 degrees of freedom
## Multiple R-squared:  0.08249, Adjusted R-squared:  0.07786
## F-statistic:  17.8 on 1 and 198 DF,  p-value: 3.731e-05

```

- The **intercept** of $\beta_0 = 44.46714$ means that a house with an **X2** value or “**house age (unit: year)**” of **0** is expected (predicted) to have a **unit price (Ping)** of **~44.47**. Since the **unit** is in “**Ping**” and each **Ping** is worth **\$10,000**, the expected value of a house that is **brand new** (0 years old) is **\$444,671.40 per Ping**. Finally, this means that the **total predicted price** of a **brand new house** is found by multiplying its **total size (Ping)** by **\$444,671.40 Taiwanese dollars**. Note: This is only true given model **Y_X2** which only models variable X2 (house age) and Y (house price) and therefore all other things are already held equal (since there is only 1 predictor variable).
- The **slope** of $\beta_1 = -0.33841$ means that for every unit change in variable **X2** or for **every 1 unit increase in house age (1 additional year)** the expected decrease of **house price** is **~0.338**. This can be interpreted to mean that **older houses** are expected to be **cheaper** or that as the **house age** increases → **house price per Ping decreases**. The precise estimate is a **decrease of \$3,384.1 per Ping** since the unit is **\$10,000 per unit of Ping**.

```

#Q1.b Continued
#Computer R-Squared
summary(Y_X2)$r.squared

## [1] 0.08249071

#Training data MSE
train_pred <- predict(Y_X2, newdata = valuation_train)
train_mse_X2 <- mean((valuation_train$Y - train_pred)^2)
train_mse_X2

## [1] 166.6527

#Test data MSE
test_pred <- predict(Y_X2, newdata = valuation_test)

```

```

test_mse_X2 <- mean((valuation_test$Y - test_pred)^2)
test_mse_X2

## [1] 188.0664

```

- The R^2 is equal to **0.08249071**. The training MSE is **166.6527**. The test MSE is **188.066**.
- c. Repeat part (b) twice, but with X2 replaced by X3 and X4 (skip the interpretation of coefficients part). Among X2, X3 and X4, which variable explains the most variation of Y on the training data? Which variable is the most useful in predicting the house price for the test data? Explain your answers.

Solution:

```

#Q1.c
# Model X3
Y_X3 <- lm(Y~X3, data = valuation_train)

train_pred_X3 <- predict(Y_X3, newdata = valuation_train)
train_mse_X3 <- mean((valuation_train$Y - train_pred_X3)^2)

test_pred_X3 <- predict(Y_X3, newdata = valuation_test)
test_mse_X3 <- mean((valuation_test$Y - test_pred_X3)^2)

# Model X4
Y_X4 <- lm(Y~X4, data = valuation_train)

train_pred_X4 <- predict(Y_X4, newdata = valuation_train)
train_mse_X4 <- mean((valuation_train$Y - train_pred_X4)^2)

test_pred_X4 <- predict(Y_X4, newdata = valuation_test)
test_mse_X4 <- mean((valuation_test$Y - test_pred_X4)^2)

summary(Y_X2)$r.squared

## [1] 0.08249071

summary(Y_X3)$r.squared

## [1] 0.5152052

summary(Y_X4)$r.squared

## [1] 0.3726554

```

- Given that the R^2 of Y_X2 is **0.08249071**, the R^2 of Y_X3 is **0.5152052**, and the R^2 of Y_X4 is **0.3726554**: X3 (“distance to the nearest MRT station) explains **the most**

variation in Y in the training data. It has the highest R^2 value (**0.5152052**), meaning it accounts for approximately **51.5%** of the variation in house price per Ping.

#Q1.c Continued

```
test_mse_X2
```

```
## [1] 188.0664
```

```
test_mse_X3
```

```
## [1] 113.7552
```

```
test_mse_X4
```

```
## [1] 135.5365
```

- **X3** is the **most useful predictor** of house price of unit area for the test data. This is because the **model using X3** has the **lowest test MSE** among the three models (Y_X2, Y_X3, Y_X4), meaning **Y_X3** makes the most accurate predictions on test (unseen) data. Note: MSE of **Y_X2 = 188.0664**, **Y_X3 = 113.7552**, **Y_X4 = 135.5365**. **MSE of X3 < X4 < X2.**
- d. Use the training data `valuation_train` to fit a multiple linear regression model of Y on X2, X3 and X4 simultaneously. Interpret all the coefficients (including the intercept) in terms of the real context. Compute R^2 , training/test MSE, and compare the results to part (b) and (c).

Solution:

```
#Q1.d
#Multiple Linear Regression of Y ~ X2+X3+X4
Y_X2X3X4 <- lm(Y~X2+X3+X4, data = valuation_train)
summary(Y_X2X3X4)

##
## Call:
## lm(formula = Y ~ X2 + X3 + X4, data = valuation_train)
##
## Residuals:
##     Min      1Q      Median      3Q      Max 
## -37.962  -5.139   -0.927    4.028   29.328 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 44.6950373  1.9558434  22.852 < 2e-16 ***
## X2          -0.3064168  0.0524983  -5.837 2.18e-08 ***
## X3          -0.0053395  0.0006038  -8.842 5.43e-16 ***
## X4           1.2502186  0.2814432   4.442 1.49e-05 ***
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

## Residual standard error: 8.455 on 196 degrees of freedom
## Multiple R-squared:  0.6143, Adjusted R-squared:  0.6084
## F-statistic: 104 on 3 and 196 DF, p-value: < 2.2e-16

```

- The β_0 represents the **expected unit price (per Ping)** for a house given **X2 = 0, X3 = 0, and X4 = 0**. In other words, given that the house is **brand new**, is **0 meters from an MRT station**, and that there are **no nearby convenience stores**... the predicted cost is **44.6950373 X \$10,000 per Ping**, which is around **\$446,950.37**. Note: It's impractical that any house would be 0 meters from an MRT station since that would mean the house is located on top of the MRT station (not possible). Therefore, this is simply a baseline prediction given the multiple linear regression model.
- The β_1 of **X2 (House Age)** is **-0.3064168**. This means that for every additional **1 year of house age**, the **expected unit price per Ping decreases by 0.3064168**. This means that **older houses are cheaper** and that for each additional year, the unit price is expected to reduce by **\$3,064 per Ping** all other variables held equal.
- The β_1 of **X3 (Distance from MRT)** is **-0.0053395**. This means that for every **1 meter increase in distance to the nearest MRT station**, the **expected unit price per Ping** decreases by **0.0053395**. This means that houses **further** from an MRT station are **cheaper**. For every additional meter from an MRT station, the expected unit price drops by around **\$53.40** all other variables held equal.
- The β_1 of **X4 (number of convenience stores in the living circle on foot)** is **1.2502186**. This means that for every additional **convenience store** within **walking distance**, the **expected unit price per ping** increases by **1.2502186**. This means that houses in an area with **more convenience** stores nearby tend to be **more valuable**. Each additional store is expected to increase the unit price by around **\$12,502.19**.

```

#Q1.d Continued
#Computing R-Squared
summary(Y_X2X3X4)$r.squared

## [1] 0.6142555

#Training MSE
train_pred <- predict(Y_X2X3X4, newdata = valuation_train)
train_mse_X2X3X4 <- mean((valuation_train$Y - train_pred)^2)

#Test MSE
test_pred <- predict(Y_X2X3X4, newdata = valuation_test)
test_mse_X2X3X4 <- mean((valuation_test$Y - test_pred)^2)

train_mse_X2X3X4

## [1] 70.06507

test_mse_X2X3X4

## [1] 100.0254

```

- The R^2 of **Y_X2X3X4 (multiple regression)** is **0.6142555**. The **training MSE** of **Y_X2X3X4** is **70.06507**. The **test MSE** of **Y_X2X3X4** is **100.0254**.

```
#Q1.d Continued
```

```
summary(Y_X2)$r.squared
## [1] 0.08249071

summary(Y_X3)$r.squared
## [1] 0.5152052

summary(Y_X4)$r.squared
## [1] 0.3726554

summary(Y_X2X3X4)$r.squared
## [1] 0.6142555

train_mse_X2
## [1] 166.6527

train_mse_X3
## [1] 88.05616

train_mse_X4
## [1] 113.9483

train_mse_X2X3X4
## [1] 70.06507

test_mse_X2
## [1] 188.0664

test_mse_X3
## [1] 113.7552

test_mse_X4
## [1] 135.5365

test_mse_X2X3X4
## [1] 100.0254
```

- The **multiple regression model (Y_X2X3X4)** has the **highest R^2 (0.6142555)**. This means that it **explains the most variation** in house price per Ping.

- The **multiple regression model ($Y_{-}X_2X_3X_4$)** has the **lowest training MSE (70.06507)**. This means that it is the best fit of the training data.
 - The **multiple regression model ($Y_{-}X_2X_3X_4$)** has the **lowest test MSE (100.0254)**. This means that it is the **most accurate at predicting unseen data**.
 - Among the **single predictor models**, **X3** performs the **best** with an R^2 of **0.5152052** and a test MSE of **113.7552**.
 - Among the **single predictor models**, **X2** is the **weakest** predictor, with the lowest R^2 and highest test MSE.
- e. Use the knnreg function in the **caret** package to fit a KNN model of Y on X_2 , X_3 and X_4 simultaneously, for each K (number of neighbors) value from 1 to 20 with increment 1. Visualize the trends of training and test MSEs in a plot where K is on the horizontal axis and MSEs are on the vertical axis. What is the optimal K and best test error? Compare the result with linear regression in part (d).

Solution:

```
#Q1.e
# Fitting KNN models for values 1:20
library(caret)

## Warning: package 'caret' was built under R version 4.4.3
## Warning: package 'ggplot2' was built under R version 4.4.3

k_seq <- 1:20

mse_seq_tr <- mse_seq_te <- NULL

for (i in seq_along(k_seq)) {
  fit <- knnreg(
    Y~X2+X3+X4,
    data = valuation_train,
    k = k_seq[i]
  )

#Training MSE
train_pred <- predict(fit, newdata = valuation_train)
mse_seq_tr[i] <- mean((valuation_train$Y - train_pred)^2)

#Test MSE
test_pred <- predict(fit, newdata = valuation_test)
mse_seq_te[i] <- mean((valuation_test$Y - test_pred)^2)

}

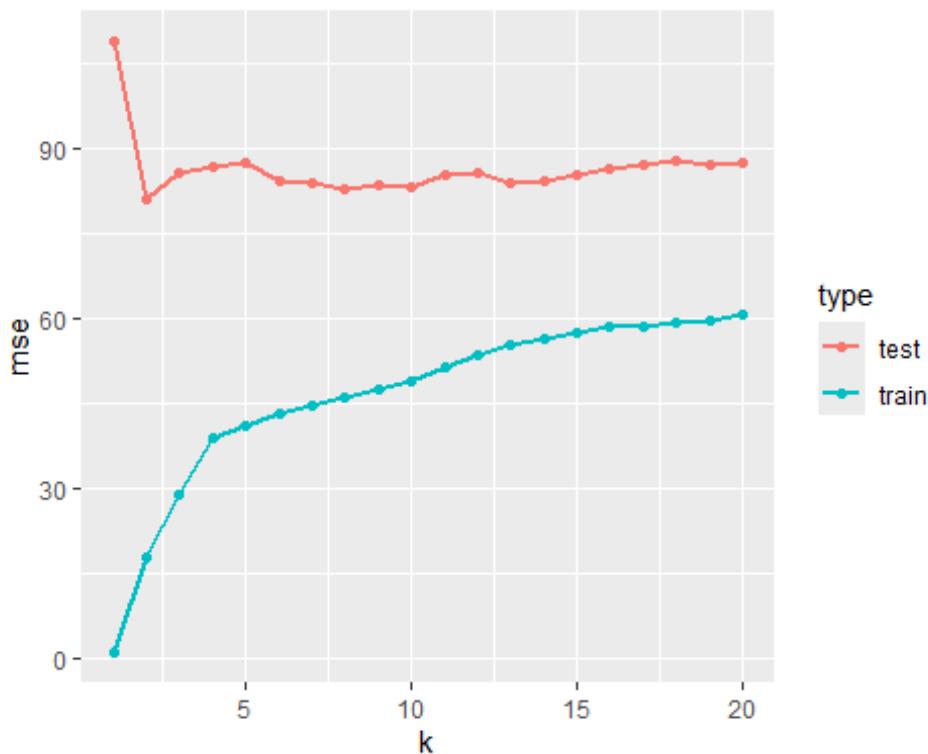
#Visualize the trends of training and test MSE
```

```

library(ggplot2)
mse_stacked <- rbind(
  data.frame(k = k_seq, mse = mse_seq_tr, type = "train"),
  data.frame(k = k_seq, mse = mse_seq_te, type = "test")
)

ggplot(mse_stacked, aes(x = k, y = mse, color = type)) +
  geom_line(linewidth = 1) +
  geom_point()

```



```

#Q1.e Continued
#Optimal K and best test error
min_test_mse <- min(mse_seq_te)
min_test_mse

## [1] 81.09504

best_k <- k_seq[which.min(mse_seq_te)]
best_k

## [1] 2

```

- The **best performing KNN model** (using X2, X3, X4) has **K = 2 neighbors** and a **test MSE** equal to **81.09504**. Therefore, **K = 2** gives the most accurate predictions on unseen data. Compared to the **multiple linear regression model** which has a **test MSE of 100.0254**, the **best KNN model (K=2)** has the **lower test MSE of 81.09504**,

which makes it a better model for the valuation data set.

Q2

The `iris` data (available in base R) gives the measurements in centimeters of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are `setosa`, `versicolor`, and `virginica`. Our goal is to predict the species using the given measurements.

To make this into a binary classification problem, we remove the `setosa` species from `iris`:

```
data(iris)
iris <- droplevels(iris[which(iris$Species != "setosa"), ])
set.seed(1)
iris[, 1:4] <- iris[, 1:4] + rnorm(400)
```

- a. Split the dataset into training and test data, using the following instructions: Put rows 1–30 and 51–80 of `iris` in the training data; put rows 31–50 and 81–100 in the test data.² Name the training and test data `iris_train` and `iris_test`, respectively. Make sure that `iris_test[34, 1]` outputs 7.500214.

Solution:

```
#Q2. a
iris_train <- iris[c(1:30, 51:80), ]
iris_test <- iris[c(31:50, 81:100), ]

iris_test[34,1]
## [1] 7.500214
```

- b. Use the training data `iris_train` to fit a logistic regression model, a LDA model and a QDA model of `Species` on all other variables. Compute the training and test errors for each model (logistic regression/LDA/QDA). Which model has the best prediction accuracy for the `iris` dataset? Explain your answer.

Solution:

```
#Q2.b
#Fitting Logistic regression model
log_fit <- glm(
  Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
  data = iris_train,
```

² Rationale behind the above split strategy: The 1–50 rows of the `iris` data frame are of species `versicolor`, whereas the 51–100 rows are of `virginica`. By combining rows 1–30 and 51–80, we end up with a balanced mix of both species in the training data. The same goes for the test data.

```

family = "binomial"
)

#fitting LDA model
library(MASS)

## Warning: package 'MASS' was built under R version 4.4.3

lda_fit <- lda(
  Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
  data = iris_train
)

#fitting QDA model
qda_fit <- qda(
  Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
  data = iris_train
)

#Log training and test error
log_train_prob <- predict(log_fit,
                           newdata = iris_train,
                           type = "response")

class_labels <- levels(iris_train$Species)
c0 <- class_labels[1]
c1 <- class_labels[2]

log_train_label <- ifelse(log_train_prob
                           > 0.5,
                           c1,
                           c0)

log_train_error <- mean(log_train_label != iris_train$Species)

log_test_prob <- predict(log_fit, newdata = iris_test,
                          type = "response")

log_test_label <- ifelse(log_test_prob
                           > 0.5,
                           c1,
                           c0)

log_test_error <- mean(log_test_label != iris_test$Species)

#LDA training and test error
lda_train_pred <- predict(lda_fit, newdata = iris_train)
lda_train_error <- mean(lda_train_pred$class != iris_train$Species)

```

```

lda_test_pred <- predict(lda_fit, newdata = iris_test)
lda_test_error <- mean(lda_test_pred$class != iris_test$Species)

#QDA training and test error
qda_train_pred <- predict(qda_fit, newdata = iris_train)
qda_class <- qda_train_pred$class
qda_train_error <- mean(qda_class != iris_train$Species)

qda_test_pred <- predict(qda_fit, newdata = iris_test)
qda_class <- qda_test_pred$class
qda_test_error <- mean(qda_class != iris_test$Species)

#Log, LDA, QDA errors
log_train_error

## [1] 0.2333333

log_test_error

## [1] 0.15

lda_train_error

## [1] 0.2333333

lda_test_error

## [1] 0.125

qda_train_error

## [1] 0.2333333

qda_test_error

## [1] 0.25

```

- The model that has the **best prediction accuracy** is the one with the **lowest test error**. This is because test error measures prediction accuracy on unseen data, which is a more true test for predictive capability. **LDA** has the best prediction accuracy for the *iris* data set. This is because **LDA achieved the lowest test error (0.125)**, which is lower than both the *logistic regression* (0.15) and *QDA* (0.25).