

Assaf Bitton

We will be predicting whether the housing price is expensive or not using the sahp dataset in the **r02pro** package.

You can run the following code to prepare the analysis.

```
library(r02pro)      #INSTALL IF NECESSARY

## Warning: package 'r02pro' was built under R version 4.4.3

library(tidyverse)  #INSTALL IF NECESSARY

## Warning: package 'ggplot2' was built under R version 4.4.3

## Warning: package 'tidyverse' was built under R version 4.4.2

## Warning: package 'dplyr' was built under R version 4.4.2

library(MASS)

## Warning: package 'MASS' was built under R version 4.4.3

my_sahp <- sahp %>%
  na.omit() %>%
  mutate(expensive = sale_price > median(sale_price)) %>%
  mutate(expensive = as.factor(expensive)) %>%
  dplyr::select(gar_car, liv_area, oa_qual, expensive)
```

Please answer the following questions.

Q1

Use the data `my_sahp` to fit four models — Logistic regression, LDA, QDA, and **standardized** KNN (with $K = 39$) — of `expensive` on all other variables. Visualize the ROC curves for them and add the AUC values to the legend. Discuss your findings. (Note: You can use the entire dataset without splitting for this analysis.)

Solution:

```
#Q1
#Logistic Regression
log_f <- glm(expensive ~.,
              data = my_sahp,
              family = "binomial")
log_p <- predict(log_f,
                  type="response")

#LDA
lda_f <- lda(expensive ~.,
              data = my_sahp)
lda_p <- predict(lda_f)$posterior[,2]
#QDA
```

```

qda_f <- qda(expensive~.,
               data = my_sahp)
qda_p <- predict(qda_f)$posterior[,2]
#Standardized KNN (K=39)
library(caret)

## Warning: package 'caret' was built under R version 4.4.3

fit_std <- preprocess(my_sahp, method = "scale")
my_sahp_std<-predict(fit_std,newdata=my_sahp)

knn_f <-knn3(expensive~.,
               data=my_sahp_std,
               k=39)
knn_p <- predict(knn_f, newdata=my_sahp_std, type="prob")[,2]
#Visualize ROC curves and add AUC to the legend
library(pROC)

## Warning: package 'pROC' was built under R version 4.4.3

log_roc <- roc(my_sahp$expensive, log_p)
log_auc <- auc(log_roc)

lda_roc <- roc(my_sahp$expensive, lda_p)
lda_auc <- auc(lda_roc)

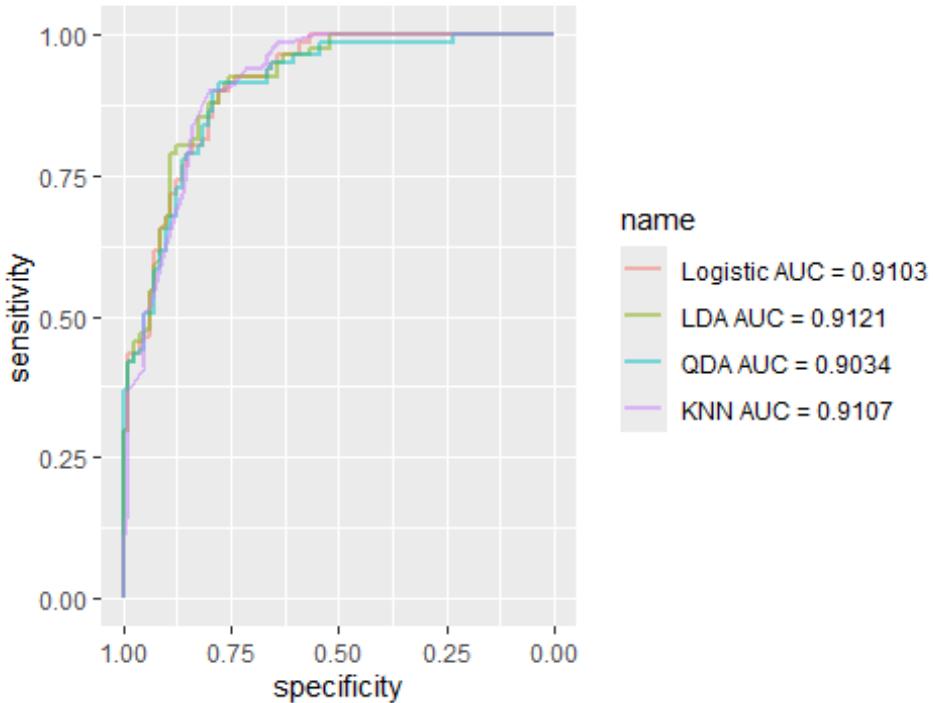
qda_roc <- roc(my_sahp$expensive, qda_p)
qda_auc <- auc(qda_roc)

knn_roc <- roc(my_sahp$expensive, knn_p)
knn_auc <- auc(knn_roc)

models_roc <- list(Logistic = log_roc,
                     LDA = lda_roc,
                     QDA = qda_roc,
                     kNN = knn_roc)
methods_auc <- paste(c("Logistic",
                       "LDA","QDA","KNN"),
                      "AUC =", 
                      round(c(log_auc, lda_auc, qda_auc, knn_auc), 4))
ggroc(models_roc, linewidth = 1, alpha = 0.5) +
  scale_color_discrete(labels = methods_auc) +
  ggtitle(paste("ROC using logistic, LDA, QDA, and KNN"))

```

ROC using logistic, LDA, QDA, and KNN



- All four models have similar predictive capabilities based on the ROC curves. Models logistic regression, LDA, QDA, and standardized KNN ($K=39$) have AUC values around 0.90 to 0.91. This indicates strong predictive capability for differentiating between expensive and non expensive homes. LDA has the highest AUC at ~0.9121, followed by KNN and logistic regression at ~0.9107 and ~0.9103 respectively, and finally QDA with the lowest AUC value of ~0.9034. This suggests that the relationship between predictors and response is modeled well by linear decision boundaries and that quadratic boundaries (QDA) does not improve performance. Overall, the models result in almost identical AUC values with no single model drastically outperforming the others.

Q2

For questions (a)–(c), use the data `my_sahp` to fit the following five models of expensive on all other variables:

- Model 1: logistic regression.
- Model 2: LDA.
- Model 3: QDA.
- Model 4: KNN with $K = 39$.
- Model 5: **standardized** KNN with $K = 39$. (use `caret::preProcess` function on the **training** sets for data standardization; see Lecture 4 slides)

For each question below, you should

- Manually split the dataset (instead of using functions like `cv.glm`). This is to ensure that each of the above models uses the same training and validation sets, so that it makes a fair comparison.
- Whenever you split the dataset (e.g., using the `sample` function as we learned in class), you should put `set.seed(1)` right above the data-splitting code to ensure reproducibility.
 - a. Use the validation set approach to divide the data into training (50%) and validation. Compute the validation classification error for each model and decide which model is the best.

Solution:

```
#Q2.a
#splitting data
n <- nrow(my_sahp)

set.seed(1)

tr_ind <- sample(n, round(n/2))
sahp_tr <- my_sahp[tr_ind,]
sahp_val <- my_sahp[-tr_ind,]

# Logistic
class_labels <- levels(my_sahp$expensive)
c0 <- class_labels[1]
c1 <- class_labels[2]

log_fit <- glm(expensive~., data = sahp_tr, family = binomial)
log_prob <- predict(log_fit, newdata = sahp_val, type = "response")
log_pred <- ifelse(log_prob > 0.5, c1, c0)

# LDA
lda_fit <- lda(expensive~., data = sahp_tr)
lda_pred <- predict(lda_fit, newdata = sahp_val)$class

# QDA
qda_fit <- qda(expensive~., data = sahp_tr)
qda_pred <- predict(qda_fit, newdata = sahp_val)$class

# KNN (K=39)
knn_fit <- knn3(expensive~., data = sahp_tr, k = 39)
knn_pred <- predict(knn_fit, newdata = sahp_val, type = "class")

# KNN (Standardized, K=39)
fit_std <- preprocess(sahp_tr, method = "scale")
sahp_tr_std <- predict(fit_std, newdata = sahp_tr)
sahp_val_std <- predict(fit_std, newdata = sahp_val)
```

```

knn_std_fit <- knn3(expensive~, data = sahp_tr_std, k = 39)
knn_std_pred <- predict(knn_std_fit,
                         newdata = sahp_val_std,
                         type = "class")

# Validation errors
cv_error_seq <- c(
  Logistic = mean(log_pred != sahp_val$expensive),
  LDA = mean(lda_pred != sahp_val$expensive),
  QDA = mean(qda_pred != sahp_val$expensive),
  KNN = mean(knn_pred != sahp_val$expensive),
  KNN_std = mean(knn_std_pred != sahp_val$expensive)
)
cv_error_seq

## Logistic      LDA      QDA      KNN      KNN_std
## 0.2345679  0.2222222  0.2222222  0.2839506  0.2098765

best_model <- names(which.min(cv_error_seq))
best_model

## [1] "KNN_std"

```

- The **standardized KNN model** has the **lowest validation error** at around 21%, making it the **best** model under the validation set approach.

b. Use LOOCV approach to compute the CV classification error for each model and decide which model is the best.

****Solution:****

```

# Q2.b
#LOOCV approach (CV classification error) for each model
log_seq <- factor(levels = class_labels)
lda_seq <- factor(levels = class_labels)
qda_seq <- factor(levels = class_labels)
knn_seq <- factor(levels = class_labels)
knnstd_seq <- factor(levels = class_labels)

for (j in 1:n){
  sahp_tr <- my_sahp[-j,]
  sahp_te <- my_sahp[j,]

  # Logistic
  log_fit <- glm(expensive~, data = sahp_tr, family = "binomial")
  log_prob <- predict(log_fit, newdata = sahp_te, type = "response")
  log_seq[j] <- ifelse(log_prob>0.5,c1,c0)

  #LDA

```

```

lda_fit <- lda(expensive~, data = sahp_tr)
lda_seq[j] <- predict(lda_fit, sahp_te)$class

#QDA
qda_fit <- qda(expensive~, data = sahp_tr)
qda_seq[j] <- predict(qda_fit, sahp_te)$class

#KNN
knn_fit <- knn3(expensive~, data = sahp_tr, k=39)
knn_seq[j] <- predict(knn_fit, sahp_te, type="class")

#standardized KNN
fit_std <- preProcess(sahp_tr, method = "scale")
sahp_tr_std <- predict(fit_std, sahp_tr)
sahp_te_std <- predict(fit_std, sahp_te)

knnstd_fit <- knn3(expensive ~., data = sahp_tr_std, k=39)
knnstd_seq[j] <- predict(knnstd_fit, sahp_te_std, type="class")
}

#LOOCV classification error for each model
cv_err_loocv <- c(
  logistic = mean(log_seq != my_sahp$expensive),
  LDA = mean(lda_seq != my_sahp$expensive),
  QDA = mean(qda_seq != my_sahp$expensive),
  KNN = mean(knn_seq != my_sahp$expensive),
  KNN_std = mean(knnstd_seq != my_sahp$expensive)
)

cv_err_loocv

## logistic      LDA      QDA      KNN      KNN_std
## 0.1790123 0.1790123 0.1851852 0.2654321 0.1913580

best_model_loocv <- names(which.min(cv_err_loocv))
best_model_loocv

## [1] "logistic"

```

- The **classification errors** for the five models are: - **Logistic regression = 0.1790123** - **LDA = 0.1790123** - **QDA = 0.1851852** - **KNN (K=39) = 0.2654321** - **Standardized KNN (K=39) = 0.1913580** - **Logistic regression** and **LDA tied for lowest LOOCV classification error (~17.9%)**. Therefore, the **logistic regression and LDA** perform the **best** under *LOOCV validation method*. QDA and standardized KNN perform slightly worse, while un-standardized KNN performs substantially worse.

- c. Use 5-fold CV approach to compute the CV classification error for each model and decide which model is the best.

Solution:

```
log_seq_fold <- factor(levels = class_labels)
lda_seq_fold <- factor(levels = class_labels)
qda_seq_fold <- factor(levels = class_labels)
knn_seq_fold <- factor(levels = class_labels)
knnstd_seq_fold <- factor(levels = class_labels)

K <- 5
set.seed(1)
fold_ind <- sample(rep_len(1:K, n))

for(k in 1:K){
  sahp_tr <- my_sahp[fold_ind != k,]
  sahp_te <- my_sahp[fold_ind == k,]

#Logistic
log_fit <- glm(expensive ~., data = sahp_tr, family=binomial)
log_prob <- predict(log_fit, newdata=sahp_te, type="response")
log_seq_fold[fold_ind==k] <- ifelse(log_prob>0.5,c1,c0)

#LDA
lda_fit <- lda(expensive~., data = sahp_tr)
lda_seq_fold[fold_ind == k] <- predict(lda_fit, sahp_te)$class

#QDA
qda_fit <- qda(expensive~., data = sahp_tr)
qda_seq_fold[fold_ind==k] <- predict(qda_fit, sahp_te)$class

#KNN
knn_fit <- knn3(expensive~., data = sahp_tr, k=39)
knn_seq_fold[fold_ind==k] <- predict(knn_fit,sahp_te,type="class")

#standardized KNN
fit_std <- preprocess(sahp_tr, method="scale")
sahp_tr_std <- predict(fit_std, sahp_tr)
sahp_te_std <- predict(fit_std, sahp_te)

knnstd_fit <- knn3(expensive~., data = sahp_tr_std,k=39)
knnstd_seq_fold[fold_ind==k] <- predict(knnstd_fit, sahp_te_std, type="class")
}

cv_err_fold <- c(
  Logistic = mean(log_seq_fold != my_sahp$expensive),
  LDA = mean(lda_seq_fold != my_sahp$expensive),
  QDA = mean(qda_seq_fold != my_sahp$expensive),
  KNN = mean(knn_seq_fold != my_sahp$expensive),
```

```

KNN_std = mean(knnstd_seq_fold != my_sahp$expensive)
)

cv_err_fold

## Logistic      LDA      QDA      KNN      KNN_std
## 0.1975309 0.2037037 0.1851852 0.2592593 0.2037037

best_model_fold <- names(which.min(cv_err_fold))
best_model_fold

## [1] "QDA"

```

The **classification errors** for the five models are: - **Logistic regression = 0.1975309** - **LDA = 0.2037037** - **QDA = 0.1851852** - **KNN (K=39) = 0.2592593** - **Standardized KNN (K=39) = 0.2037037** - The **QDA model** has the **lowest** 5 fold CV error at **0.1851852**. Therefore, QDA performs the **best** among the five models under this evaluation.

Q3

ISLRv2 Chapter 4 Q4 (Page 189). To receive full credits, you need to show your work.

Solution: Q4. (A) Given we have **1 predictor variable** ($p = 1$) that is uniform $[0,1]$: If we use observations within 10% of the range, **we expect to use 10% of the observations for each prediction**. So on average, we will use **10% of observations for each prediction**.

- (B) Given we have **2 predictor variables** ($p = 2$) that is uniform $[0,1]$: If we use observations within 10% of the range for each predictor, **we form a square that has features: $[x_1 \pm 0.05] \times [x_2 \pm 0.05]$** whose area is **$0.1 \times 0.1 = 0.01$** , therefore the **fraction of observations used is: 1%**.
- (C) Given we have **100 predictor variables** ($p = 100$) that is uniform $[0,1]$: If we use observations within 10% of the range for each predictor, **we form a 100 dimensional hyper cube with features**: Each side has a length of 0.10 and therefore **$(0.1)^{100}$ represents the 100 dimensional hyper volume of the hyper cube**. Therefore, the value is incredible small (1×100^{-100}) . This means **the fraction of observations used for the prediction is 10^{-100}** .
- (D) As the dimension **p increases**, the **volume** of the hyper shape becomes **$(0.1)^p$** . This **shrinks exponentially** toward 0. Therefore, in *low dimensions* there **are nearby points**, but in *high dimensions* there **are almost no points nearby**. In high dimensional space, KNN uses extremely distant (non-local) points which leads to bad predictions.
- (E) Given we want to create a p -dimensional hypercube centered around the test observation that contains 10% of the training observations on average: Let the side of the length be l :

- The **volume of the hypercube is l^p**
- The **fraction of points** it contains must be: $l^p = 0.10$
- Therefore, the **side length is: $l = 0.10^{1/p}$**
- Now, we plug in for **$p = 1, 2, \text{ and } 100$** .
- **For $p = 1$, the length = $0.10^{1/1} = 0.10$**
- **For $p = 2$, the length = $0.10^{1/2} = 0.316$**
- **For $p = 100$, the length = $0.10^{1/100} = 0.97725$** This means that as **p dimensions increase**, the hyper cube must **extend closer to the full range of data** and therefore the predictions using KNN **no longer becomes local and essentially useless**.